

Last Time:

□ Intro to HRL

lecture 2

HRI, FALL '25

This Time:

□ Sequential decision-making

□ MDPs

□ Policies, Values, Bellman Eqn.

Andrea Bojcsy

What is sequential decision-making?

Sequential decision-making involves making a sequence of decisions over time, where each decision affects future outcomes and decisions.


↳ different from one-shot decisions which are a self-contained single decision (e.g. img. classification)

Why do we care?


Sequential decision-making is everywhere - playing games, planning a trip or career - and it's present in interaction!

In HRI, sequential decision-making will form the "mathematical backbone" of how we model people, robots, and their interaction:


- robots must plan ahead, considering long-term consequences
- humans are sequential decision-makers too: must model them well
- interaction is an ongoing exchange where agents influence each other



How will R fly home?




How will H get to their car?



How should R hand H the mug?

Q Why is it hard?

A1 decision tree grows exponentially in time horizon.
 ↳ too many possibilities to compute



$$|A| = 3$$

Starting from ~~the~~ state
 there are

$$|A|^{T+1}$$


sequences of decisions

$$= 3^4 = 81 \text{ sequences}$$


with time-horizon of 3! \square

⇒ Naively planning quickly becomes impossible!

A2 outcomes of taking actions can be stochastic
 ↳ same action doesn't always have the same result.




e.g. R chooses to go straight,
 but wind might push it
 left or right.



⇒ Even with same action sequence,
 outcomes may not be the same!

A3 Hard to assign credit to the right past action(s)
 ↳ delayed rewards make it unclear which action was responsible.



✓ successfully home, but was it all good...?
 ✗ fail to get home, but was it all bad...?

⇒ If rewards come late, hard to tell which
 actions were good / bad!

Markov Decision Processes (MDPs)

MDPs are a mathematical model for sequential decision-making in a fully observable, stochastic, discrete-time environment with a Markovian transition model.

$$M = \langle S, A, T, r \rangle$$

$$A = \{\leftarrow, \uparrow, \rightarrow\}$$


- states: $s \in S$

state space

- actions: $a \in A$

action space

- transition function: $T: S \times A \rightarrow S$



$$T(s, a, s') \text{ or}$$

$$T(s=(1,3), a=\uparrow, s'=(2,3)) = 0.8$$

$$P(s' | s, a) \text{ or}$$

$$P(s'=(2,3) | s=(1,3), a=\uparrow) = 0.8$$

"is distributed as"
 $s' \sim P(\cdot | s, a)$

$$s' \sim P(\cdot | s=(1,3), a=\uparrow) = \begin{cases} (2,3) \rightarrow 0.8 \\ (1,2) \rightarrow 0.1 \\ (1,4) \rightarrow 0.1 \end{cases}$$

- reward function: $r: S \times A \times S \rightarrow \mathbb{R}$

$$\textcircled{R} \quad r(s=(2,3), a=\rightarrow, s'=(2,4)) = -1$$

$$r(s'=(2,4)) = -1$$

$$r(s'=(3,4)) = +1$$

$$r(s, a) \text{ or } r(s, a, s')$$

or

② What is "Markov" about this?

The transition function is where the markov part comes into play!

- Markov Property: the future is independent of the past, given the present.

This means that action outcomes depend only on the current state, not history!



↓

$$P(s_{t+1} | s_t, a_t, s_{t-1}, a_{t-1}, \dots, s_0) = P(s_{t+1} | s_t, a_t) \text{ makes planning easier}$$

"Get the Gem"

"Get the gem after visiting A"

markov holds!



Markov does not hold!

BUT, if $s = (x, y, A_visited)$... yes!

! If a problem is Markov depends on how we define the state

Policies

The object we ultimately want to solve for is the policy: OR, knowing what is the best decision to make @ any state.

$$\pi : S \rightarrow A$$

"deterministic, $\pi(s)$ "

$$\pi : S \rightarrow \Delta A$$

"stochastic: $\pi(a|s)$ "

probability simplex, "distribution over actions"

$\pi : S \rightarrow A$

ex. policy (not → very good one tho...)

| | | | | |
|---|---|---|---|----|
| 3 | → | → | → | +1 |
| 2 | → | | → | -1 |
| 1 | → | → | → | → |

1 2 3 4

Evaluate the quality of a policy π by the expected cumulative reward

induced by the policy. ⚡ we will get to the "expected" bit in a bit ...

- ① "roll out" or simulate π from each $s_0 \in S$ to get a trajectory sequence $\{s_0, \pi(s_0), s_1, \pi(s_1), \dots\}$

- ② Evaluate the utility of the rollout, called cumulative reward:

$$R(s_0, s_1, s_2, \dots) \triangleq r(s_0) + r(s_1) + \dots = \sum_{t=0}^{\infty} r(s_t) \quad \begin{cases} \text{or} \\ \text{r}(s_t, a_t, s_{t+1}) \end{cases}$$

$\pi_1 : S \rightarrow A$

| | | | | |
|---|---|---|---|----|
| 3 | → | → | → | +1 |
| 2 | → | | → | -1 |
| 1 | → | → | → | → |

1 2 3 4

$\pi_2 : S \rightarrow A$

| | | | | |
|---|---|---|---|----|
| 3 | → | → | → | +1 |
| 2 | ↑ | | ↑ | -1 |
| 1 | ↑ | → | ↑ | ← |

1 2 3 4

↳ better policy b/c more reward!

• an optimal policy ($\pi^* : S \rightarrow A$) yields the highest cumulative reward

- ! But what counts as a "good" policy depends on the reward function!

Sparse vs. Dense Rewards


Problem: If the only reward is a +1 at the goal (and -1 @ lava), many policies look equally good because there is no penalty for inefficiency! In other words, sparse rewards often underspecify what we "actually" want the agent to do!

(*) Underspecified rewards is a HUGE problem in HRL and AI Alignment!!!

Solution: Dense rewards - frequent feedback (good / bad) to the agent over time, instead of just rewarding outcomes at the end.

- + helps guide the agent to behave in the way we actually want.
 - + specify not just what the agent should do (i.e. goal), but how to do it.
 - + faster agent learning
 - HARD TO SPECIFY + PROVIDE → see later lectures on this!

Example: "Living" rewards are a kind of dense reward that incentivizes efficiency by penalizing necessary steps.



Discount Factor

- denoted by $\gamma \in [0, 1] \Rightarrow$ sometimes MDP tuple is written as $\langle S, A, T, r, \gamma \rangle$
 - describes an agent's preference for current rewards over future ones
when $\gamma < 1$. If $\gamma \equiv 1$, our agent wants max reward over all tsteps.

$$R(s_0, s_1, s_2, \dots) = \gamma^0 r(s_0) + \gamma^1 r(s_1) + \gamma^2 r(s_2) + \dots = \sum_{t=0}^{\infty} \gamma^t r(s_t)$$

discounted cumulative reward

* discounting is a good model of human & animal preferences

↳ See: "models of Temporal Discounting 1937-2000"
by Till Grüne-Yanoff.

→ connections w/ psych + econ.

Ultimately, we are searching for a policy $\pi^*: S \rightarrow A$ that maximizes

the expected discounted cumulative reward:

$$\pi^* = \underset{\pi}{\operatorname{argmax}}$$

"return the policy π which maximizes..."

$$E \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \right]$$

trade off shorter vs.
longer horizon
rewards

$\tau \sim p(\tau)$ rollouts are stochastic so take expectation
let $\tau := (s_0, a_0, s_1, a_1, \dots)$ over all trajectories under
the policy π !

$$p(\tau) = p(s_0) \prod_{t=0}^{\infty} p(s_{t+1} | s_t, a_t)$$

⊕ note: this is for
deterministic π ! for
stochastic we have
 $= p(s_0) \prod_t p(a_t | s_t) p(r_{t+1} | s_t)$